


Enterprise RAG

Zero-hallucination and 100% evidence backed system for enterprise documents




A hybrid semantic+graph search outperforms traditional retrieval by up to 50%



	Structural Context	Multimodal Assets	Relational View	Agentic retrieval
Complication	Enterprise documents contain highly structured components like logical paragraphs, complex tables, and explicit headers	Corporate intelligence is dense with non-textual assets, including financial charts, trend graphs, and process diagrams	Enterprise analysis relies heavily on synthesizing cross-document connections and hierarchies	Raw user queries are often poorly framed, lack dense domain keywords, and fail to match target vector embeddings directly.
Situation	Naive solutions shred document layouts, causing incomplete retrieval and wasting billable hours re-verifying data	Standard vector search is blind to visual elements, dropping non-textual data and creating blind spots in quantitative reporting	Vector search suffers from topological blindness, only surfacing isolated snippets while failing to detect other relationships	Single-pass vector search suffers from input dependency; bad phrasing misses critical documents completely.
Resolution	Layout-aware parsing to isolate document boundaries, preserving complete semantic units and achieving 100% content retention	Multimodal vision models and embeddings to make figures searchable, securing 100% visual coverage	Dual-index architecture pairing semantic vector search with a dedicated graph database for 50% better accuracy on multi-hop queries	Agentic reasoning loop that decomposes inputs, enhances query, runs hybrid search, evaluates sufficiency and retries as needed

Architecture Diagram

Modular, secure architecture built to process complex data accurately and block errors automatically



Implementation roadmap

The project will be implemented in 4 distinct phases, with room for review and changes.

